

# Joaquin Mendoza

530-536-6018 | [Email](#) | [LinkedIn](#) | [GitHub](#) | [portfolio.joaquinmendoza.social](#)

## EDUCATION

### San Diego State University

Aug. 2021 – Dec. 2024

*B.S. in Computer Science*

*San Diego, CA*

- **GPA:** 3.78/4.00 Cumulative
- **Relevant Courses:** Machine Learning, Data Science, Neural Networks, Database Theory, Software Engineering
- **Extracurriculars:** Waterski Club (Vice President), AI Club (Project Lead), Pickleball Club (D1 Team)

## EXPERIENCE

### Data Engineer Intern

May 2024 – Aug. 2024

*GoFundMe Inc.*

*San Diego, CA (FT)*

- Reduced data model complexity by 43%, accelerating financial reporting for over 4,200 charities by revamping a Snowflake-based KPI workflow.
- Unified data from 50,000+ transactions by engineering a mapping diagram for fee types, aligning data across Snowflake transaction tables, Looker dashboards, and Confluence for simplified fee revenue analyses.
- Piped data from 30 Zuora charge plans into Salesforce, refining the ETL process for accurate revenue forecasting by developing a dbt model for Stripe revenue aggregation.

### Software Engineer Intern

May 2023 – May 2024

*Real Frequency – Launch Agency*

*Spokane, WA (Remote-FT)*

- Reduced annual labor costs by \$50k by designing and implementing a Flask-based data reporting app.
- Increased business operations efficiency by 50% through a 240-hour reduction in monthly data entry time using scripted data workflows.
- Boosted customer acquisition by 100% through leveraging business intelligence to optimize engagement with 500+ monthly active users.

### Software Engineer Assistant

Dec. 2022 – Dec. 2024

*IT Division, San Diego State University*

*San Diego, CA (PT)*

- Automated microservice deployments for 7 clusters using Helm charts and GitLab CI/CD pipelines, enhancing support for machine learning and parallelized workloads.
- Optimized high-performance computing (HPC) resource allocation for 15 STEM courses and research labs through Kubernetes pod spawning logic.
- Containerized model training and inference for 11 machine learning jobs on the National Research Platform (NRP) hypercluster: Nautilus.

## PROJECTS

### Hubble Finances (Personal) | *Vue.js, MySQL, Bokeh*

May. 2024 – Present

- Scaled a cloud-native personal finance app to render interactive visualizations on top of Plaid's production API.
- Secured end-to-end encryption for all stored and managed personally-identifiable information (PII) to comply with state and federal data privacy regulations.

### Real Insights (Real Frequency) | *JQuery, Flask, Plotly*

Aug. 2023 – Jan. 2024

- Developed a business intelligence (BI) app, transforming raw user data into dynamic dashboards, using Jinja2.
- Automated 10 hours of weekly spreadsheet reporting, alleviating bottlenecks in continuous client analysis.
- Produced 200+ monthly interactive reports delivering actionable insights to 37 sports organizations at zero-cost.

### Spotify Atlas (CS 574) | *Scikit-Learn, Matplotlib, Pandas*

Feb. 2024

- Trained a song recommendation engine on a 114,000-song dataset, enabling personalized music discovery.
- Applied PCA analysis on the Spotify dataset to address dimensionality overhead and reduce memory constraints.

## TECHNICAL SKILLS

**Languages:** Python, SQL, C, JavaScript

**Data Tools:** dbt, Snowflake, MySQL, Hex, Looker

**Dev Tools:** Docker, Kubernetes, Git, New Relic

**Libraries:** Bokeh, Plotly, Matplotlib, Flask