Joaquin Mendoza

530-536-6018 | [Email](#) | [LinkedIn](#) | [GitHub](#) | portfolio.joaquinmendoza.social

EDUCATION

San Diego State University — Aug. 2021 – Dec. 2024
B.S. in Computer Science — *San Diego, CA*
- **GPA**: 3.78/4.00 Cumulative
- **Relevant Courses**: Machine Learning, Data Science, Neural Networks, Database Theory, Software Engineering
- **Extracurriculars**: Waterski Club (Vice President), AI Club (Project Lead), Pickleball Club (D1 Team)

EXPERIENCE

Business Intelligence Analyst Intern — May 2024 – Aug. 2024
GoFundMe Inc. — *San Diego, CA (FT)*
- Unified data from 50,000+ transactions by engineering a mapping diagram for fee types, aligning data across Snowflake transaction tables, Looker dashboards, and Confluence to simplify sales revenue analytics.
- Piped data from 30 Zuora charge plans into Salesforce, refining the ETL process for accurate revenue forecasting by developing a dbt model for Stripe revenue aggregation.
- Reduced data model complexity by 43%, accelerating data-driven financial reporting across 4,200 charities, while consulting with business operations and post-sales leaders to satisfy cross-functional business needs.

Software Engineer Intern — May 2023 – May 2024
Real Frequency – Launch Agency — *Spokane, WA (Remote-FT)*
- Achieved 100% growth in customer acquisition by leveraging product analytics to present data-driven insights to C-suite leadership, driving strategic changes in customer engagement and product strategy.
- Increased business operations efficiency by 50% through a 240-hour reduction in monthly data entry time using scripted data workflows.
- Reduced annual labor costs by $50k through the design and implementation of a custom Flask-based data reporting tool.

Software Engineer Assistant — Dec. 2022 – Dec. 2024
IT Division, San Diego State University — *San Diego, CA (PT)*
- Automated microservice deployments for 7 clusters using Helm charts and GitLab CI/CD pipelines, enhancing support for machine learning and parallelized workloads.
- Optimized high-performance computing (HPC) resource allocation for 15 STEM courses and research labs through Kubernetes pod spawning logic.
- Containerized model training and inference for 11 machine learning jobs on the National Research Platform (NRP) hypercluster: Nautilus.

PROJECTS

Hubble Finances (Personal) | *Vue.js, mySQL, Bokeh* — May 2024 – Present
- Scaled a cloud-native personal finance app to deliver visualizations derived from advanced analytics using Plaid's REST API.
- Secured end-to-end encryption for all stored and managed personally-identifiable information (PII) to comply with state and federal data privacy regulations.

Real Insights (Real Frequency) | *JQuery, Flask, Plotly* — Aug. 2023 – Jan. 2024
- Developed a business intelligence (BI) app, transforming raw user data into dynamic dashboards using Jinja2.
- Produced 200+ monthly interactive reports delivering key data analytics to 37 sports organizations at zero-cost.

Spotify Atlas (CS 574) | *Scikit-Learn, Matplotlib, Pandas* — Feb. 2024
- Trained a song recommendation engine on a 114,000-song dataset, enabling personalized music discovery.
- Applied PCA analysis on the Spotify dataset to address dimensionality overhead and reduce memory constraints.

TECHNICAL SKILLS

Languages: Python, SQL, C, JavaScript
Data Tools: dbt, Snowflake, MySQL, Hex, Looker, Census, Excel
Dev Tools: Docker, Kubernetes, Git, New Relic
Libraries: Bokeh, Plotly, Matplotlib, Flask